

DiRe: Topology-faithful dimensionality reduction at scale

Alexander Kolpakov^a and Igor Rivin^{b,1}

^aUniversity of Austin, Austin, TX, USA; ^bDepartment of Mathematics, Temple University, Philadelphia, PA, USA

This manuscript was compiled on April 26, 2026

We revisit the evaluation of dimensionality-reduction algorithms in light of the observation that metrics based on k -nearest-neighbor preservation reward faithful reproduction of sampling noise as well as manifold structure. We propose a direct replacement: *topology error* on noisy-manifold stress tests, measuring deviation from known first Betti numbers in the 2-D embedding via *riper*-computed persistence diagrams. Using this metric alongside classification accuracy as objectives in a multi-objective NSGA-II search, we tune the DiRe force-directed embedding method and compare against GPU-accelerated cuML UMAP on 11 OpenML datasets and on a 723-thousand-paper arXiv embedding corpus. On 7 of 11 OpenML datasets, tuned DiRe strictly Pareto-dominates cuML UMAP on both axes; on every dataset, a DiRe configuration exists that recovers exact first Betti numbers on both stress manifolds while UMAP never drops below a topology error of 5. On the arXiv corpus, DiRe preserves 3–4× more topological structure (Betti-curve DTW distance) than UMAP at comparable wall-clock (20 s vs 32 s for 723K points on a single H100-class GPU). The resulting hyperparameter preset is published as part of the open-source DiRe library.

dimensionality reduction | topological data analysis | visualization | manifold learning | GPU algorithms

Dimensionality reduction methods such as t-SNE (1) and UMAP (2) are the standard tools for two-dimensional visualisation of high-dimensional datasets, from single-cell expression profiles to language model embeddings. Both methods optimise a local-neighborhood-preservation objective: the embedding is good to the extent that each point's k nearest neighbors in the original space remain its k nearest neighbors in the embedding. This design emphasises cluster structure and short-range relationships, and both methods are effective at it.

The local objective is, however, known to come at the expense of global and topological structure. UMAP in particular produces embeddings whose 2-D layouts introduce gaps, folds, and spurious cycles that are not in the underlying data (3, 4). These artefacts are visually striking and scientifically misleading: two papers appearing in separate UMAP islands need not be more dissimilar than two papers in the same island, yet published biological UMAPs are routinely read as though they were.

We make three contributions. First, we argue that the local-neighborhood metric rewards reproduction of sampling noise as well as manifold structure. We demonstrate this by sweeping noise level on manifolds with known topology and showing that UMAP's β_1 estimate tracks the inflating noisy-sample β_1 , while a force-directed alternative (DiRe, (5)) stays near the theoretical value (Figure 1).

Second, we propose *topology error* — the absolute deviation of β_1 in the 2-D embedding from the known β_1 of a fixed pair of stress manifolds — as a direct, scale-invariant measurement

of topological faithfulness. Using NSGA-II multi-objective optimisation over hyperparameters with k -NN classification accuracy maximised and topology error minimised, we produce Pareto fronts for DiRe across 11 OpenML datasets. We observe strict Pareto dominance of cuML UMAP on 7 of 11; on every dataset there exist DiRe configurations that reach topology error 0 (Figure 2).

Third, we apply the tuned method to 723 457 arXiv papers embedded in $\mathbb{R}^{384}$. DiRe preserves the reference Betti curves of the high-dimensional corpus 3–4× more faithfully than UMAP (Figure 4) in comparable wall-clock.

Method

The DiRe embedding. DiRe (5, 6) is a three-stage method. (i) A k -NN graph of the input point cloud is computed on the GPU via cuVS (7). (ii) An initial low-dimensional embedding is produced by one of: Principal Component Analysis, Laplacian Eigenmaps via LOBPCG on the normalised adjacency (8), diffusion maps (9), or Johnson–Lindenstrauss random projection (10). (iii) A force-directed layout refines the embedding using an attractive/repulsive potential

$$\Phi_{\text{att}}(d) = \log(1 + a d^{2b}), \quad \Phi_{\text{rep}}(d) = -\log\left(\frac{a d^{2b}}{1 + a d^{2b}}\right),$$

with (a, b) fit from *spread* and *min_dist* parameters, but evaluated with a *bounded* force kernel and accumulated in mixed precision for efficient GPU execution. The full pipeline is available in *dire-rapids*, with *init='spectral'* the effective default from the studies reported below.

Significance Statement

Dimensionality reduction methods such as UMAP and t-SNE are central tools for visualising high-dimensional data, but their local-neighborhood objectives can preserve sampling noise while distorting global topology. We show that standard local metrics reward this noise memorisation: top-performing embeddings invent cycles and disconnected islands absent from the data. We introduce a topology-faithfulness benchmark based on noisy manifolds with known homology, tune DiRe against it, and find Pareto-optimal configurations that match or beat GPU-accelerated UMAP on classification while recovering exact first Betti numbers on stress tests. On 723K arXiv paper embeddings, DiRe preserves 3–4× more topological structure than UMAP at comparable wall-clock.

A.K. and I.R. designed research, performed research, contributed new reagents/analytic tools, analyzed data, and wrote the paper.

The authors declare no competing interests.

¹To whom correspondence should be addressed. E-mail: rivintemple.edu

σ	noisy-sample	DiRe-pca	DiRe-spec	cuML UMAP
0.01	2.0	2.1	2.0	2.2
0.02	2.0	2.0	1.9	1.5
0.05	2.0	1.9	2.0	1.9
0.10	2.1	1.9	1.8	1.9
0.20	38.4	2.5	3.0	7.7

Table 1. β_1 significant-bar count on figure-8 (theoretical $\beta_1 = 2$). Each entry is mean over 10 seeds at 1000 points.

Why local-neighborhood metrics can be misleading. The standard k -NN preservation metric rewards an embedding for matching the *exact* set of k nearest neighbors of each point. On a noisy sampled manifold, which specific k points are closest depends both on the manifold geometry and on the noise draw. An embedding that crinkles itself to reproduce both the signal and the noise scores higher than one that preserves only the signal. The noise sweep (Figure 1) shows this quantitatively on figure-8.

Topology error. We define *topology error* as a single scalar:

$$\text{TE} = \sum_{M \in \mathcal{M}} |\beta_1^{\text{embed}}(M) - \beta_1^{\text{true}}(M)| \quad [1]$$

over a fixed stress-test set $\mathcal{M}$. In this work $\mathcal{M} = \{\text{figure-8 at } \sigma = 0.2, \text{ torus}^2 \text{ at } \sigma = 0.05\}$, both with theoretical $\beta_1 = 2$. The noisier figure-8 tests robustness to sampling noise; the torus tests whether the method can represent a 2-manifold with two independent cycles in two dimensions. β_1^{embed} is measured by running `ripser` (11) on the embedding, converting the persistence diagram to Betti curves, and counting significant bars (persistence ≥ 0.3 of maximum persistence).

Multi-objective hyperparameter search. For each benchmark dataset we run an NSGA-II study (12) with two objectives: maximise 2-D k -NN classification accuracy on the dataset, minimise the topology error of Equation 1. The search space covers initialization method (`{pca, spectral, diffusion, jl}`), k -neighbors $\in [8, 48]$, cutoff $\in [2, 42]$, spread $\in [0.5, 4.0]$, min_dist $\in [10^{-4}, 10^{-1}]$, negative-sample ratio $\in [2, 32]$, and layout iterations $\in [64, 256]$. We run 150–200 trials per dataset with a population of 25. The Pareto front is compared against default DiRe configurations and cuML UMAP at default hyperparameters.

Results

Noise-robustness discriminator. Table 1 reports β_1 bar counts on noisy figure-8 samples (1000 points, 10 seeds per σ). The naive noisy-sample β_1 climbs from 2.0 at $\sigma = 0.01$ to 38.4 at $\sigma = 0.2$. DiRe-pca rises only to 2.5; cuML UMAP rises to 7.7.

Pareto dominance on 11 OpenML datasets. Table 2 summarises the outcome of 150-trial NSGA-II searches on 11 OpenML datasets. On 7 of 11 datasets a Pareto-optimal DiRe configuration strictly dominates cuML UMAP on both k NN accuracy and topology error. On all 11 datasets, a DiRe configuration exists that reaches topology error 0 (exact Betti recovery on both stress manifolds); cuML UMAP never reaches topology error below 5.

figure-8 ($\beta_1 = 2$): method response to noise




Fig. 1. As sampling noise grows, the noisy point-cloud’s β_1 count inflates sharply (from 2 at $\sigma = 0.01$ to ≈ 38 at $\sigma = 0.2$). UMAP’s embedding partially tracks this inflation ($\beta_1 \approx 7.7$ at $\sigma = 0.2$); DiRe’s embedding stays close to the theoretical value ($\beta_1 \approx 2.5$ at $\sigma = 0.2$). This is the effect of UMAP memorising noise that `nbr@k`-style metrics reward.

Across the four Pareto studies run on different datasets, the Pareto-optimal hyperparameters converge on a consistent region: `init='spectral'`, `spread` $\in [2.0, 3.9]$ (vs default 1.0), `max_iter_layout` $\in [125, 242]$ (vs default 128), `n_neighbors` $\in [10, 30]$. We publish this region as a preset, `dire_rapids.presets.TOPOLOGY_TUNED`, available with `from dire_rapids import TOPOLOGY_TUNED`; `DiRePyTorch(**TOPOLOGY_TUNED)`.

arXiv paper embeddings. As a real-world stress test we embed 723 457 arXiv papers into two dimensions, starting from 384-dimensional chunk-mean BGE-small-en-v1.5 embeddings (13) of the abstracts and full texts. Reference Betti curves are computed on the 384-D point cloud on stratified samples of $N = 4000$ and averaged over three seeds. We sweep the k -neighbors parameter for both methods to rule out the possibility that the advantage we report is a single- k artefact.

Island-ness. As a sharper diagnosis of how UMAP’s topology is wrong, we measure features of the 0-dimensional persistence diagram on the same 2-D layouts (Table 5, Figure 3). UMAP’s longest H_0 bar is consistently 2–5 $\times$ DiRe’s, and the ratio of the top 5 bars to the median bar is consistently 2–5 $\times$ higher for UMAP. Both metrics say that UMAP concentrates topological mass into a few very-long bars — the signature of breaking a continuous manifold into well-separated islands. DiRe’s distribution is broader, closer to the reference point cloud’s. The arXiv category taxonomy is genuinely continuous (math.AG continuously touches math.SG, math.DG, math.AT, etc.); UMAP’s embedding presents it as disjoint pieces, DiRe’s preserves the tendril structure.

Discussion

Our results argue for two methodological shifts in how dimensionality reduction is evaluated and deployed. First, local-neighborhood preservation metrics are *necessary but not sufficient*: they can be gamed by methods that faithfully reproduce noise, and they rank those methods above alternatives that

Dataset	UMAP k NN	UMAP TE	best k NN (TE)	best TE (k NN)	UMAP dominated by
mfeat-factors	0.893	5	0.915 (7)	0.907 (0)	3/4
satimage	0.883	5	0.876 (2)	0.874 (0)	0/2
pendigits	0.987	5	0.985 (1)	0.984 (0)	0/2
isolet	0.810	5	0.842 (11)	0.833 (0)	2/3
HAR	0.918	5	0.936 (6)	0.925 (0)	3/4
letter	0.875	5	0.895 (13)	0.872 (0)	3/7
magic	0.766	5	0.762 (7)	0.748 (0)	0/8
MNIST	0.970	5	0.961 (1)	0.960 (0)	0/2
Fashion-MNIST	0.782	5	0.798 (2)	0.792 (0)	3/3
connect-4	0.648	5	0.663 (1)	0.652 (0)	2/2
covertypes	0.695	5	0.727 (4)	0.715 (0)	3/3

Table 2. NSGA-II Pareto search for DiRe versus cuML UMAP on 11 OpenML datasets. Objectives: maximise k NN accuracy in 2-D, minimise topology error (Equation 1). Bold entries indicate a strict improvement over UMAP on that axis. UMAP dominated by reports the number of Pareto trials that strictly dominate cuML UMAP on both axes simultaneously. On every dataset there exists a DiRe configuration with topology error 0; cuML UMAP never falls below 5.




Fig. 2. Pareto front from the 200-trial NSGA-II study on covertypes ($N = 581\,012$, $D = 54$). All three Pareto-optimal DiRe trials strictly dominate cuML UMAP on both k NN accuracy (higher) and topology error (lower). The best-topology trial achieves topology error 0 at k NN = 0.715, a +2pp absolute k NN improvement over cuML UMAP.

produce globally more faithful embeddings. Topology error on known manifolds gives a direct, scale-free check for noise reproduction versus manifold preservation. Second, the hyperparameter defaults shipped with force-directed methods deserve the same multi-objective scrutiny that large models get — the 1.0 spread default in DiRe turns out to be off by a factor of 3.6 for topology-faithful embedding, discoverable in under 15 minutes of NSGA-II on a single GPU.

Limitations. Ripser’s $\text{maxdim}=2$ cost grows as $O(N^3)$, which confined our topology objective to β_0 and β_1 on $N = 1000$ stress samples. Scaling the topology objective to higher-dimensional homology at larger N requires GPU-native persistent-homology infrastructure, which to our knowledge does not yet exist at research-grade reliability for dimensions ≥ 2 . Our present 2-D/3-D embedding comparisons do not probe β_k for $k \geq 3$ in the original data.

Reproducibility. All experiments are reproducible

k	method	DTW(β_0) ↓	DTW(β_1) ↓
8	DiRe	2149 ± 472	423 ± 52
8	cuML UMAP	6943 ± 1045	408 ± 70
16	DiRe	3901 ± 506	193 ± 16
16	cuML UMAP	10429 ± 466	818 ± 328
32	DiRe	5632 ± 49	192 ± 13
32	cuML UMAP	11212 ± 313	1084 ± 320
64	DiRe	7663 ± 653	275 ± 41
64	cuML UMAP	11978 ± 109	508 ± 8
128	DiRe	5389 ± 550	224 ± 63
128	cuML UMAP	12125 ± 124	664 ± 158

Table 3. Topology preservation on the 723 457-paper arXiv corpus, swept over k -neighbors. Lower Betti-curve DTW is better. DiRe wins DTW(β_0) at every k by 2–3×; DiRe wins DTW(β_1) at every k except $k = 8$ (where both are close). Means ± std over three seeds on stratified $N = 4000$ samples.

	DiRe	cuML UMAP	—
Trustworthiness	0.796	0.869	
k NN-preservation@15	0.74	0.82	
Wall-clock (GH200, 723K pts)	20.3 s	32.4 s	1.6×

Table 4. Local-neighborhood and wall-clock metrics at $k = 16$. As the noise-robustness discriminator argues, UMAP’s local-metric advantage here is at least in part the noise-memorisation effect that local metrics reward. DiRe is 1.6× faster on the same GPU.

via `benchmarking/bench_topology_pareto.py` in the `dire-rapids` open-source package. The preset is in `dire-rapids.presets.TOPOLOGY_TUNED`. arXiv-corpus-specific code, including the BGE embedding pipeline and stratified Betti evaluation, is in the companion `dire-rapids-arxiv` repository.

Materials and Methods

Hardware and software. All benchmarks run on a single NVIDIA GH200 GPU under `conda env rapids-26.04` (Python 3.14, PyTorch 2.11, RAPIDS 26.04). `diRe` is implemented in `dire-rapids` v0.3.0. UMAP baseline: cuML 26.04 with default hyperparameters ($k = 16$). TSNE baseline: cuML 26.04 with default perplexity. Topology: ripser 0.6.14 via `ripser.py`.

k	layout	longest H_0 bar $\downarrow$	top-5 / median ratio $\downarrow$
16	reference (384-D)	0.70	2.4
16	DiRe	0.13	17.3
16	cuML UMAP	0.27	71.2
32	DiRe	0.09	23.6
32	cuML UMAP	0.31	74.1
128	DiRe	0.09	32.2
128	cuML UMAP	0.39	146.5

Table 5. Island-ness: persistence-diagram features of the 2-D layouts at selected k . Longest H_0 bar (lower = more continuous, like reference); top-5-over-median ratio (lower = persistence spread across many bars rather than concentrated into a few). UMAP is 2–5 $\times$ more "island-y" than DiRe at every k .




Fig. 3. Longest H_0 persistence bar (lower = more continuous) vs. k -neighbors on the arXiv corpus 2-D layouts. Dashed line: reference persistence of the 384-D point cloud (0.70; far above both methods, because the reference manifold is one large continuum with many slow-growing bars rather than a few dominant islands). UMAP sits 2–3 $\times$ above DiRe at every k . The shape of the deviation (long leading bar) is the island-creation pathology, quantitatively.

Preprocessing. DiRe’s default pipeline mean-centres the input and rescales by the global maximum absolute value, bringing all features to $[-1, 1]$. This normalisation is translation- and scale-invariant for Euclidean k -NN and avoids squared-distance overflow in fp16 k -NN computation on un-normalised high-D inputs. UMAP is fed the same normalised inputs for fair comparison.

Topology-error metric. Stress manifolds: figure-8 (a lemniscate, $\beta_1 = 2$) sampled at $\sigma = 0.2$ Gaussian noise, and a standard torus ($R = 2$, $r = 1$, $\beta_1 = 2$) at $\sigma = 0.05$. Both sampled with 1000 points. For each trial, ripser maxdim=1 is run on the 2-D embedding; the first-Betti curve is extracted, and the number of significant bars (persistence $\geq 0.3 \times$ max persistence) is recorded. Topology error for a trial is $|b_1^{\text{fig-8}} - 2| + |b_1^{\text{torus}} - 2|$.

Pareto search. optuna 4.x NSGA-II sampler, population 25, seed 0. 150–200 trials per dataset. Search space in the Method section. Each trial fits DiRe on the target dataset, computes k NN accuracy with `KNeighborsClassifier(15)` on a fixed 70/30 train/test split, and computes topology error on the stress manifolds. Missing-value rows in OpenML datasets are dropped.

arXiv corpus. 723 457 arXiv papers embedded with BAAI/bge-small-en-v1.5 (13), chunk-level embeddings mean-pooled per paper to a 384-D vector. Reference Betti curves: ripser maxdim=1 on

stratified samples of $N \in \{1000, 4000\}$, averaged over three seeds. Embedding Betti curves: ripser on the 2-D layout at the same stratified sample indices. DTW distance: `fastdtw` on the Betti curve time series.

Data Availability

Code and scripts needed to reproduce the experiments are available in the open-source `dire-rapids` repository, including `benchmarking/bench_topology_pareto.py`. arXiv-corpus-specific code, including the BGE embedding pipeline and stratified Betti evaluation, is available in the companion `dire-rapids-arxiv` repository.

ACKNOWLEDGMENTS. The benchmarks reported in this paper were carried out on an NVIDIA GH200 kindly provided by Lambda Labs (<https://lambda.ai>); we thank Lambda for the extended access. This work was also supported in part by the Google Cloud Research Award GCP19980904.

1. L van der Maaten, G Hinton, Visualizing data using t-sne. *J. Mach. Learn. Res.* **9**, 2579–2605 (2008).
2. L McInnes, J Healy, J Melville, UMAP: Uniform manifold approximation and projection for dimension reduction (arXiv:1802.03426) (2018).
3. T Chari, L Pachter, The specious art of single-cell genomics. *PLOS Comput. Biol.* **19**, e1011288 (2023).
4. T Chari, L Pachter, The specious art of single-cell genomics (extended) (bioRxiv preprint) (2023).
5. A Kolpakov, I Rivin, DiRe-JAX: a JAX-based dimensionality reduction algorithm (Journal of Open Source Software, forthcoming) (2025).
6. A Kolpakov, I Rivin, `dire-rapids`: PyTorch + RAPIDS accelerated dimensionality reduction (<https://github.com/sashakolpakov/dire-rapids>) (2026).
7. NVIDIA RAPIDS Team, cuVS: CUDA vector search (<https://github.com/rapidsai/cuvs>) (2024).
8. AV Knyazev, Toward the optimal preconditioned eigensolver: Locally optimal block preconditioned conjugate gradient method. *SIAM J. Sci. Comput.* **23**, 517–541 (2001).
9. RR Coifman, S Lafon, Diffusion maps. *Appl. Comput. Harmon. Analysis* **21**, 5–30 (2006).
10. WB Johnson, J Lindenstrauss, Extensions of lipschitz mappings into a hilbert space. *Contemp. Math.* **26**, 189–206 (1984).
11. U Bauer, Ripser: efficient computation of Vietoris–Rips persistence barcodes. *J. Appl. Comput. Topol.* **5**, 391–423 (2021).
12. K Deb, A Pratap, S Agarwal, T Meyarivan, A fast and elitist multiobjective genetic algorithm: NSGA-II. *IEEE Transactions on Evol. Comput.* **6**, 182–197 (2002).
13. BAAI, BGE: BAAI general embedding (small, english, v1.5) (<https://huggingface.co/BAAI/bge-small-en-v1.5>) (2023).




Fig. 4. arXiv-paper 2-D layout, DiRe (left) vs cuML UMAP (right), coloured by primary arXiv category. Both methods recover category structure; UMAP produces tighter but more fragmented islands with visible gaps in long-range category similarity (e.g. math.DG vs math.SG). DiRe's layout is smoother and preserves the inter-category topology of the reference 384-D point cloud more faithfully.